TIPOFF · PHASE 1 · V2 · OCTOBER 2026

Where the model stands

Prediction-market signals → merchant decisions.

Backtested on 296 real events across the NBA, NHL, NFL and MLB (two seasons each).

Everything in this deck is generated from the backtest outputs (pre-registration sha256 8625f39966bc…). Demand is measured with search interest, not real orders.

THE ANSWER

Our signal makes merchants money. The market's number is the edge only in the NFL.

Proven in backtest

$366

extra profit per event vs doing nothing

(store selling 20 units/day; +16% of the maximum possible)

Partly

+37% vs +18%

NFL: our signal vs a smart merchant's gut feel

Overall it's only +16% vs +13%, not a reliable edge

Where the money is

+70%

of the opportunity captured on **"stop"** calls after eliminations

Big repricings: +35%. Ordinary good news: +14%.

WHAT WE SET OUT TO PROVE

The chain, and how far along each link is

1	**Spike → real signal** 780 odds moves → 546 kept. Noise (single wallets, reverting prints, one-venue moves) filtered without hindsight.	Proven
2	**Signal → plain-English insight** Every kept move gets a cause: series clinch, elimination, playoff win/loss, rival's game, trade or injury.	Proven
3	**Demand moves after the signal** Fans' merch searches follow the signal's direction over the next two weeks (NBA, NFL, MLB).	Proven
4	**Merchant acts at T=0** Five merchant policies replayed on every event. No real store has acted on a signal yet.	Simulated
5	**Merchant makes more money** Yes vs doing nothing ($366/event). Vs a smart merchant: only clearly in the NFL.	Simulated

"Proven" = shown in historical data with out-of-sample tests. "Simulated" = replayed with a model of merchant behaviour; needs real store orders to confirm.

HOW THE MODEL WORKS

Five steps, all using only information available at the moment the signal fires

1 · Sense

Hourly title odds for every team, Polymarket + Kalshi (as a cross-check), 4 leagues × 2 seasons

780 candidate moves

2 · Filter

Move must hold 6 hours, Kalshi must agree, enough trades, not driven by one wallet

546 kept

3 · Classify

Which game or news caused it? Clinch, win, loss, elimination, title, rival's game, off-court news

8 event types

4 · Forecast

Expected 14-day demand for this type of event, adjusted by how far the odds moved. Trained on the other season only

no lookahead

5 · Act

Order size (or stop) at T=0, plus a playbook per merchant type

one decision per event

QUESTION 1

Can we tell real signals from noise? Yes

If a move is real, its **size** should predict how much demand changes. If it's noise, it shouldn't.

League	Moves → kept	Kept: size predicts demand (t)	Rejected (t)
NFL	138 → 115	+10.5	+2.4
MLB	217 → 126	+6.3	+5.8
NBA	147 → 126	+2.8	-0.4
NHL	278 → 179	+0.9	+1.6

t above ~2 = a real relationship. Kept moves pass in the NFL, MLB and NBA. Rejected moves are ~0 in the NBA (as they should be), but in thin MLB markets our filter also threw out some real moves: thresholds need tuning per market. NHL: too little demand data.

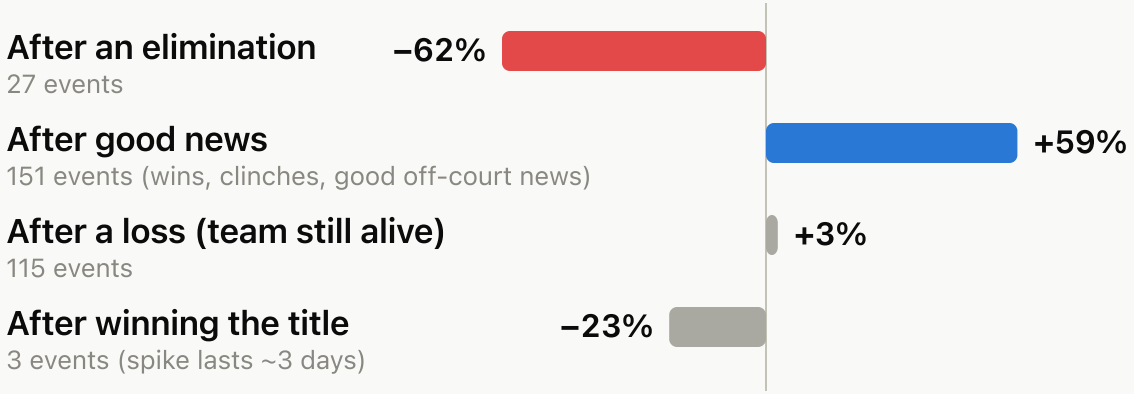
Example of noise caught

October 2024, before the season: one wallet pushed the Pacers' title odds from 1.6% to 6%. No game, no news. A naive "odds spiked" alert would have told Indiana merchants to stock up seven months early. Our filter rejected it, and fan demand never moved.

QUESTION 2

Does demand actually move after a signal?

Yes, and the type of event matters more than the arrow



Demand over the next 14 days vs the week before the signal.

Direction test (does demand follow the odds?): p = <0.001 NFL, <0.001 MLB, 0.002 NBA.

The trap a raw odds alert falls into: a playoff loss drops a team's title odds, but demand **doesn't** fall while the team is still playing. Only eliminations collapse demand. That's why we classify every move before acting.

QUESTION 3 · YOUR THREE SCENARIOS

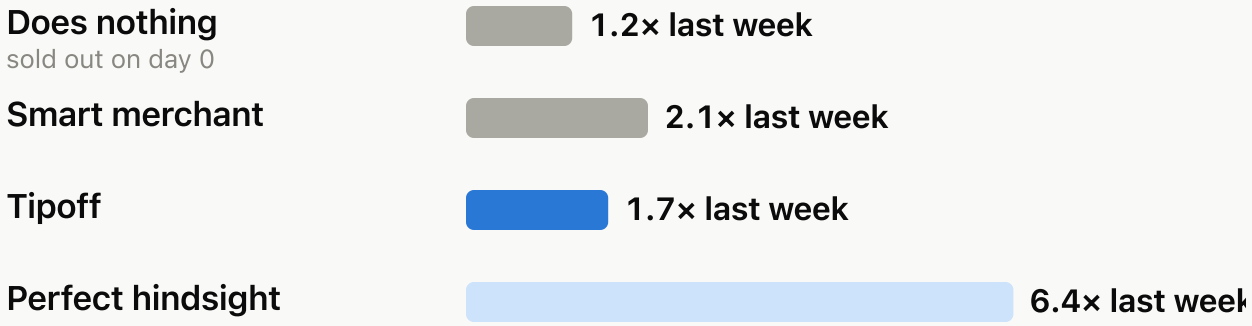
If a merchant got our signal at T=0, would they make more money? Yes

Merchant	What they do at T=0	Units sold vs last week	Extra profit per event	Share of the max possible
1 · Does nothing	Re-orders 2 weeks at last week's pace	−7%	baseline	0%
2 · Smart merchant	Knows the headline, orders what that kind of news usually means	+2%	$307	+13%
3 · Tipoff	Our event type + how far the odds moved	+3%	**$366**	**+16%**
Perfect hindsight	Knew exactly what demand would be	+25%	$2,347	100%

296 events. One order at T=0 covering 14 days, delivered instantly, no top-ups. Store sells 20 units/day; $30 retail, $12 cost, 30% recovered on leftovers. Tipoff's 90% confidence range: +6% to +25% of the max, entirely above zero. Demand over these windows ran +25% vs the week before; "does nothing" ran out of stock on the big days and over-stocked after eliminations.

WORKED EXAMPLE · LUKA DONČIĆ TRADED TO THE LAKERS, FEB 2, 2025

The signal fired at 8am. Demand came in at 6.4× the previous week's pace.



How much stock each merchant ordered at T=0, as a multiple of what they sold the week before (scaled to 14 days).

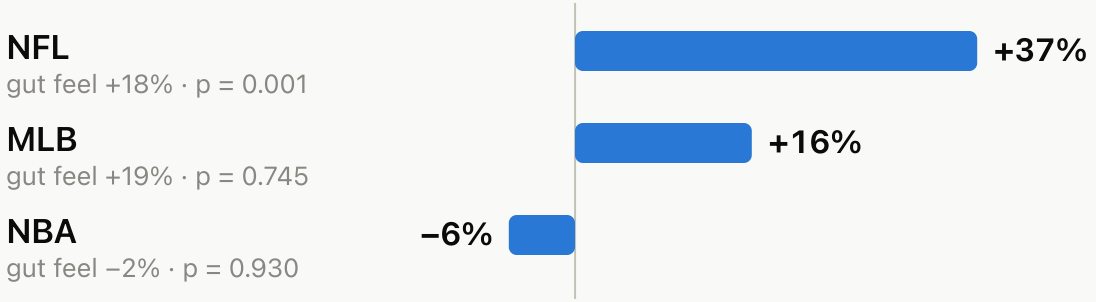
Merchant	Profit (14 days)
Does nothing	$6,241
Smart merchant	$10,701
Tipoff	$8,360
Perfect hindsight	$32,184

Honest read: everyone under-ordered. Tipoff beat doing nothing by $2,119 but ordered less than the smart merchant. Rare shocks this big are the hardest thing to size: there are too few of them to learn from.

QUESTION 4 · THE REFRAME

Is the market's number better than a smart merchant's gut?

Only in the NFL



Tipoff's share of the max possible profit, by league. NHL: too few events with usable data (US merch searches are sparse).

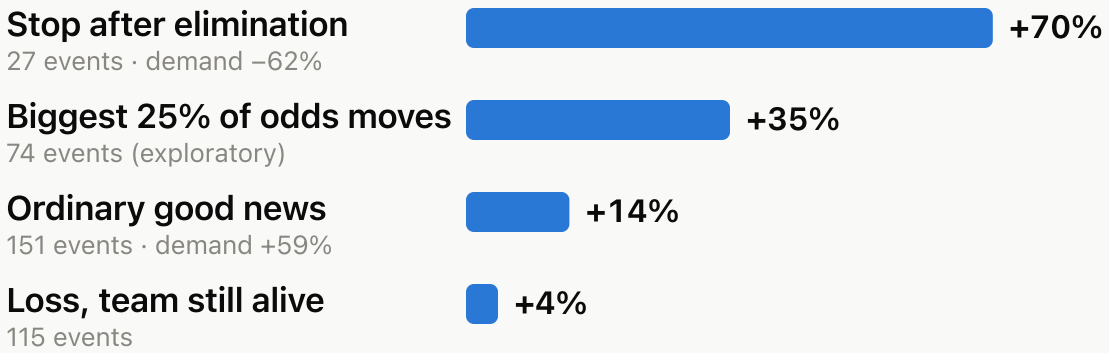
Overall: Tipoff +16% vs smart merchant +13% (p = 0.210): not a reliable edge.

NFL: Tipoff +37% vs +18% (p = 0.001), ahead in all 6 robustness variants (+11 to +22 points).

Why it might be the NFL: 17 games, so every result moves the season; one decision a week; the biggest merch market. This was not a pre-registered hypothesis, so it needs confirming on the live 2026 season.

WHERE THE MONEY IS

"Stop" and "big moves" pay most. Ordinary good news pays less.



Tipoff's share of the max possible profit, by type of event.

Stop: when a team is eliminated, demand falls about 60% within two weeks. A merchant who simply stops restocking that night captures most of the opportunity. That's the easiest win to automate and to sell.

Big moves: the market is best at telling you *which* news is big enough to act on.

Ordinary good news: bump sizes vary too much from event to event to size precisely from search data.

INTEGRITY

How we kept ourselves from fooling ourselves

Locked the test before seeing results

Design written and fingerprinted (sha256 8625f39966…) before the first v2 run.

No peeking ahead

Signals use only prices up to the moment they fire; every forecast is trained on the other season.

Two fixes, both disclosed

A model blow-up on title wins (original spec: Tipoff −2%) and game moves mislabelled as news. Before/after shown on the site.

Honest uncertainty

Confidence ranges resample whole team-seasons, because one team's playoff signals are related.

Placebo test

Shuffling the odds numbers mostly erases Tipoff's edge (+13% vs +16%, p = 0.065): borderline, not conclusive.

Robustness

Results re-run with 7- and 21-day windows, different margins, pooled training. Overall conclusions hold; NFL edge holds in every variant.

WHAT'S NOT PROVEN YET

The gaps, in order of importance

1. **Search interest is not sales.** Every demand number is Google search volume for team merch. Real Shopify orders could move every result, up or down. **Biggest gap.**
2. **No real merchant has acted on a signal.** The merchant behaviour is simulated.
3. **Delivery delays were ignored** (your simplification). Our v1 test with real lead times: signals help slow-ordering stores most and fast print-on-demand stores least.
4. **The NFL edge was found, not predicted.** It needs an out-of-sample confirmation.
5. **Thin markets.** NHL demand data is too sparse to score; MLB's noise filter is too strict.
6. **Small sample of rare shocks.** Trades and season-ending injuries (the Luka type) are too few to size well.

WHAT IT MEANS FOR THE STARTUP

The reframe is half right, and that's enough to build on

The market as a detector: works. It tells you something happened, which way it cuts, and whether it's big.

The market as a ruler for order size: works in the NFL, not (yet) in the NBA or MLB.

So the product is an act/stop radar:

- Stop: pause ads, cancel orders, mark down the night a season ends
- Act: alert only on big repricings
- Size precisely where proven (NFL first); elsewhere give a direction and a range

Price on profit, not sales. Signal-following merchants sell about +11% more units, but some of that comes with leftover stock. A cut of sales would look good while costing the merchant money.

Best customers: stores that commit early. With real lead times (v1), signals added +9% profit for 14-day-lead licensed sellers and about nothing for print-on-demand.

Bigger prize later: long-lead decisions (imports, tariffs) where big repricings are the norm.

NEXT 30 DAYS

Turn "simulated" into "proven"

1. **Week 1:** lock a forward test on the live 2026 NFL season, with the same pre-registration discipline. Log every signal and the order each merchant type would place.
2. **Week 2:** sign 3–5 design-partner stores (NFL fan merch first), pull their order history, and re-run this exact test on real orders.
3. **Week 3:** add series and player markets; tune the noise filter per market liquidity.
4. **Week 4:** ship the act/stop radar MVP to partners; measure the actions they take and the dollars saved.

What we need

- Design-partner store orders (anonymised Shopify exports)
- Your Polymarket parquet backtester (full trade history)
- A decision: NFL-first?
- An Anthropic API key for live news classification